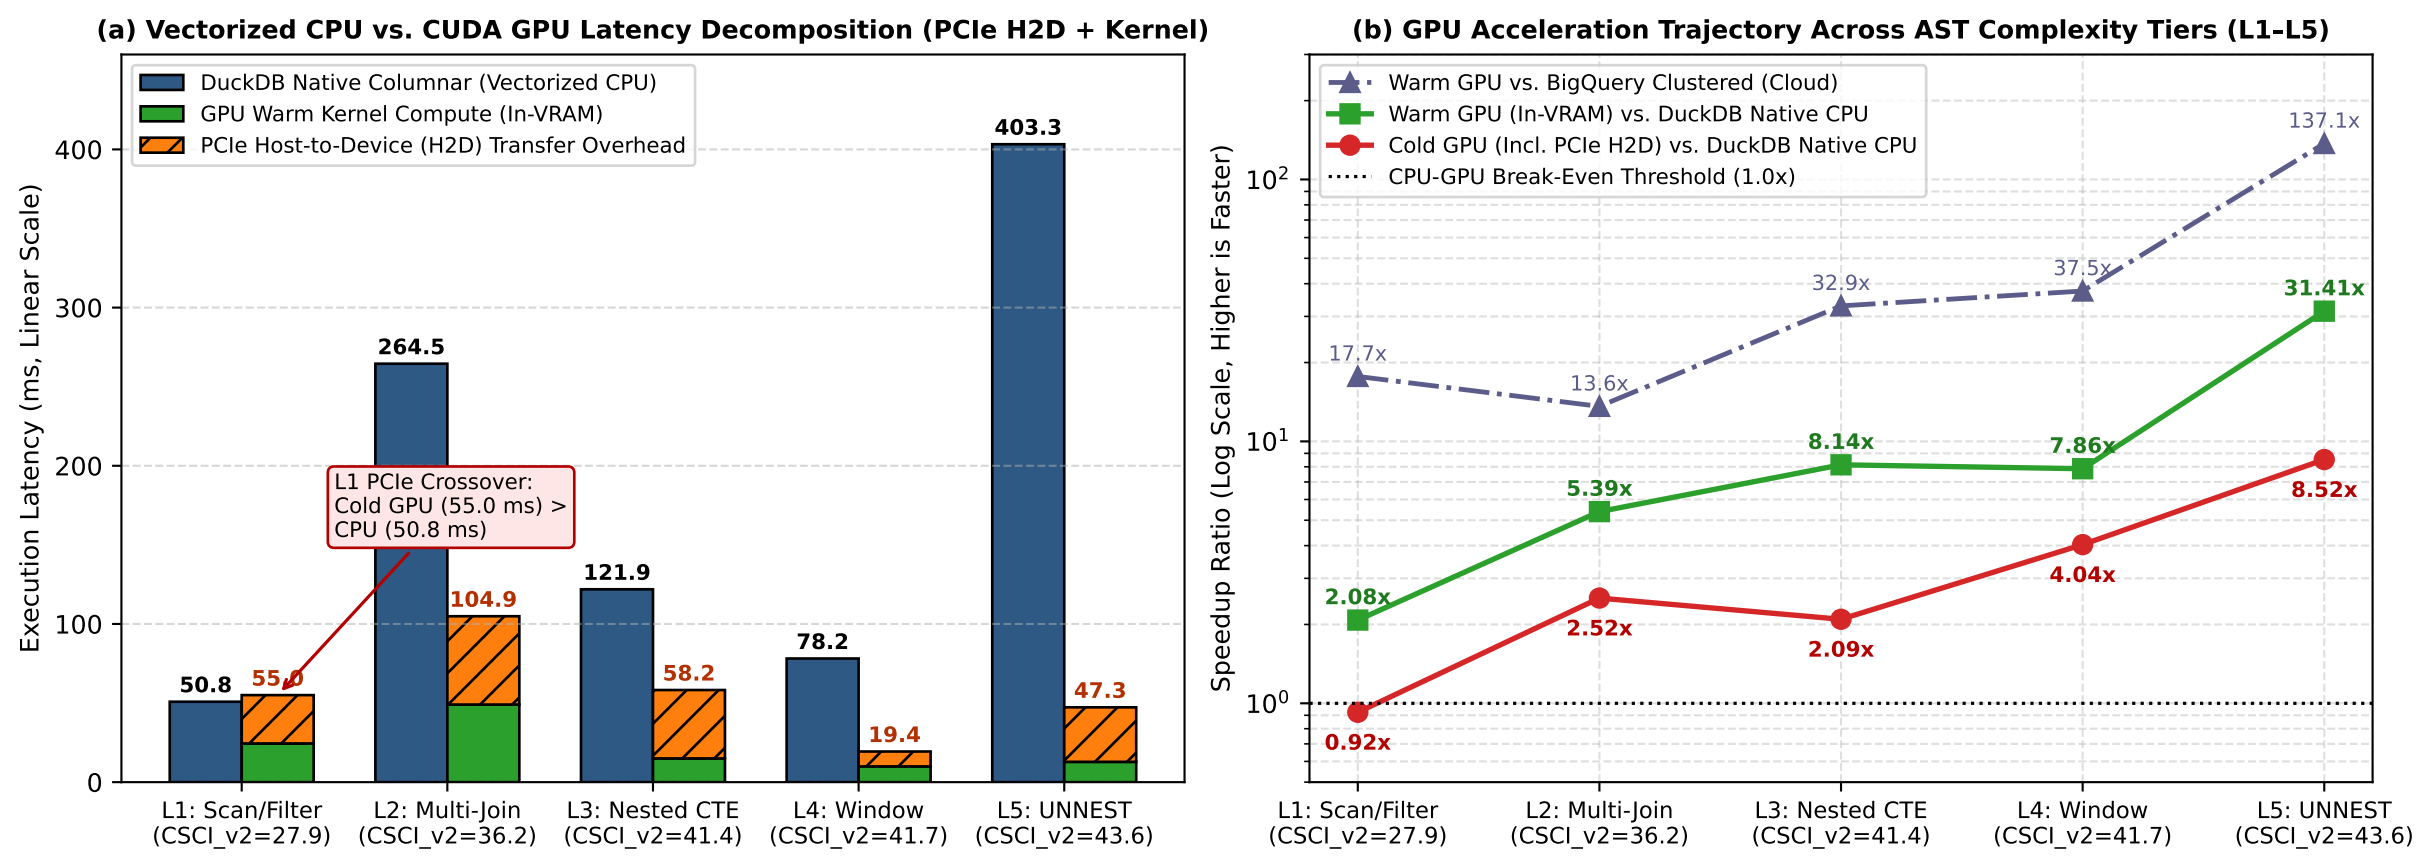
(a) Vectorized CPU vs. CUDA GPU Latency Decomposition (PCIe H2D + Kernel)
DuckDB Native Columnar (Vectorized CPU)
GPU Warm Kernel Compute (In-VRAM)
PCIe Host-to-Device (H2D) Transfer Overhead
Execution Latency (ms, Linear Scale)
0
100
200
300
400
50.8
55.0
264.5
104.9
121.9
58.2
78.2
19.4
403.3
47.3
L1 PCIe Crossover:
Cold GPU (55.0 ms) >
CPU (50.8 ms)
L1: Scan/Filter
(CSCI_v2=27.9)
L2: Multi-Join
(CSCI_v2=36.2)
L3: Nested CTE
(CSCI_v2=41.4)
L4: Window
(CSCI_v2=41.7)
L5: UNNEST
(CSCI_v2=43.6)
(b) GPU Acceleration Trajectory Across AST Complexity Tiers (L1–L5)
Warm GPU vs. BigQuery Clustered (Cloud)
Warm GPU (In-VRAM) vs. DuckDB Native CPU
Cold GPU (Incl. PCIe H2D) vs. DuckDB Native CPU
CPU-GPU Break-Even Threshold (1.0x)
Speedup Ratio (Log Scale, Higher is Faster)
10^0
10^1
10^2
17.7x
13.6x
32.9x
37.5x
137.1x
2.08x
5.39x
8.14x
7.86x
31.41x
0.92x
2.52x
2.09x
4.04x
8.52x
L1: Scan/Filter
(CSCI_v2=27.9)
L2: Multi-Join
(CSCI_v2=36.2)
L3: Nested CTE
(CSCI_v2=41.4)
L4: Window
(CSCI_v2=41.7)
L5: UNNEST
(CSCI_v2=43.6)